

How the AI Job Finder works

A plain guide to using the board, and to what actually happens when you upload a resume.

ON THIS PAGE

1. [Reading the board](#)
2. [Uploading a resume](#)
3. [What happens on a run](#)
4. [How scoring works](#)
5. [How location works](#)
6. [Suggested employers](#)
7. [What is verified, what is not](#)
8. [Where your data goes](#)
9. [When something goes wrong](#)

1 Reading the board

About eighty employers, sorted by whether you can act on them today.

The board opens with no filters applied, so you see everything at once. Cards are grouped by timing first and by match score second, which means the roles you can apply to this week sit at the top.

Nothing on the default board assumes anything about you. School-specific features stay out of the way until a resume asks for them: upload one showing a University of Utah student and a career fair strip and an alumni filter appear, because the board tracks which of these employers recruit on that campus.

-  **Live roles right now.** A real opening was confirmed open on the date in the footer. Apply this week.
-  **Suggested for your city.** Only appears after you upload a resume from outside Salt Lake. These come from Claude's knowledge and are not verified.
-  **Opens later this fall.** The program is real and dated, but applications are not open yet. Set a reminder.
-  **Watch list.** The company hires, and their board is live, but nothing entry level was posted on the check date.
-  **Wrong cohort.** The timing does not fit the graduation window. This group is hidden when it is empty.

The number on each card

That circle is a match score from 60 to 100. On the default board it rates fit for a new grad targeting data, AI and fintech work around Salt Lake City. After you upload a resume it rates fit for that person instead. Higher means a stronger case, not a better company.

Filters

Filters combine, so picking a role and a city narrows to cards matching both. Picking a second chip in the same row swaps rather than stacks. Any filter that would return nothing is hidden, so you can never land on an empty board. Search covers company names, cities, the explanation text and job titles.

2 Uploading a resume

Five steps, about a minute, usually under twenty cents.

1. Click **Upload a resume**.
2. Paste an Anthropic API key. Get one at console.anthropic.com. It is saved in your browser so you only do this once.
3. Pick a model and effort level, or keep the defaults.
4. Drop in a PDF, a Word file or a text file. You can also paste the text.
5. Click **Tailor the board** and watch it think.

Claude's reasoning streams into the panel while it works, and a counter tracks how many companies it has written. Stop cancels at any point and leaves the board untouched.

What it costs

MODEL AND EFFORT	PER RUN	TIME	WHEN TO PICK IT
Sonnet 5, medium	about \$0.15	about a minute	The default, and the right answer almost always.
Sonnet 5, low	under \$0.10	well under a minute	A quick pass when you just want the shape of it.
Opus 5, high	\$0.50 to \$1.50	several minutes	An unusual resume or a career change.
Fable 5.1	\$1.00 to \$3.00	slowest	Only when Opus 5 got something wrong.

Effort is the bigger lever, not the model. Thinking is billed as output, so a high-effort run on a big model spends most of its money reasoning rather than writing. Opus 5 at high effort was only a third of the way through the board after four minutes in testing, which is why the defaults are Sonnet 5 at medium.


The toast at the end reports what the run actually cost.

Saved profiles

Every finished run is saved as a profile and a dropdown appears in the header. Switch between people freely, or pick **Original board** to go back to the untouched version. Profiles live in that one browser and never sync anywhere.

3 What happens on a run

One request, sent straight from your browser to Anthropic. There is no server in the middle.



One round trip. Nothing is uploaded to any server belonging to this project, because there isn't one.

Why the answer is always the right shape

The request pins Claude to a JSON schema, so it cannot reply with prose or invent extra fields. The page then treats that answer as untrusted anyway: scores outside 60 to 100 are clamped, unknown categories are dropped, companies that are not on the board are ignored, duplicates are removed, and anything shown as text is escaped. A reply covering fewer than 80% of the companies is rejected rather than half applied.

4 How scoring works

A score answers one question: how strong is this person's case for an entry-level hire here?

SCORE	MEANING
95 to 100	Bullseye. Field, city, tools and hiring pipeline all line up.
85 to 94	Strong. A clear case with one soft spot.
75 to 84	Solid. Worth applying, not a standout.
65 to 74	Stretch. Needs a good cover letter or a referral.

SCORE	MEANING
60 to 64	Weak. Listed so you know it was considered and why it lost.

Scores are deliberately spread. A board where everything scores in the high eighties tells you nothing, so the rules require using the full range and being explicit when a role is a reach.

Experience is a ceiling, not a factor

A posting demanding three years cannot be a great match for someone graduating next term, however well the rest lines up. Each listed opening was checked against the live posting for what it actually requires, and anything wanting two years or more is tagged on the card. The page then caps the score itself rather than asking the model nicely: three years beyond your experience caps an employer at 70, two years caps it at 80, and if even the easiest opening there is out of reach the company moves to the watch list.

One reachable role rescues a company. An employer listing a new-grad programme alongside a senior role is judged on the new-grad one, because that is the door you can actually walk through.

Badges and tiers are worked out here, not by Claude

Claude supplies one score per employer and a single yes-or-no on whether its stack overlaps tools named on the resume. Everything else follows automatically. The location badge comes from where the employer sits relative to the candidate. Fintech and AI-native come from the researched board, because they describe the company and do not change with who is applying. Tier one is simply the eight highest scores.

This is deliberate. Asking a model to decide eighty tiers and badge sets costs real money in reasoning and produces answers that drift between runs. Deriving them is free, instant and identical every time.

Timing is separate from fit

Geography never changes whether a role is live. If a company has a real opening today, it stays in the live group even for someone across the country. The score drops instead, and the explanation says the role is on site. Timing only moves when graduation dates genuinely conflict, such as a program that starts before the person finishes school.

5 How location works

The board was built around Salt Lake City, so every location is re-read relative to whoever uploaded a resume.

LABEL	WHAT IT MEANS FOR THIS PERSON
Local	Headquarters or a major hub in their metro.
Has an office	A smaller office there.

LABEL	WHAT IT MEANS FOR THIS PERSON
Remote	Hires this kind of role remotely from where they live.
Needs relocation	Would require a move. Hidden when nothing qualifies.

The filter labels rename themselves to the person's city, so an Austin candidate sees "Austin Area" and "Has Austin office". A remote-only candidate sees neither, because those categories would be empty.

The Salt Lake exception. If the resume shows a home base anywhere on the Wasatch Front, or says the person is moving to Salt Lake, the board stays exactly as it is. No relabeling and no suggested employers. The researched board is already correct for those candidates, so rewriting it would only lose information. Two things are gated separately: the career fair strip, which follows the school rather than the city and appears only for a University of Utah student, and suggested employers, which appear for anyone whose field falls outside what this board hires for, Salt Lake resident or not.

6 Suggested employers

The escape hatch for when this particular board is the wrong board.

Sixty-odd hand-picked employers are only useful if they are employers you could actually work for. Two things break that. You might live somewhere else, in which case the Utah companies are unreachable. Or you might work in a different field, in which case they are irrelevant even from across town: nearly all of them hire data, analytics or software people.

Either way, Claude adds employers in your own city and your own field: up to fifteen, drawn from its own knowledge, mixing large employers that run early-career programs with notable local firms. They appear as dashed purple cards under a heading naming the city. Only someone who is both near Salt Lake and in one of the board's fields gets no suggestions, because for them the researched list is already the right answer.

When your field is the mismatch, the suggestions move to the top of the page and the board drops beneath them. A nurse in Salt Lake City should see hospitals first, not nineteen software companies with genuinely open roles she cannot apply for. The board says so in the group heading rather than quietly reordering.

They are handled differently from the researched board in three ways, all deliberate:

- They are labeled **not verified** on every card.
- Their button runs a web search for the company's careers page instead of linking to a job posting that may not exist.
- Each carries a confidence level, and lower-confidence entries say so.

The rules tell Claude to name a hiring program only when it is certain of the exact name, and otherwise to describe what to search for. Fewer good suggestions beat a padded list.

7 What is verified, what is not

The single most important distinction on the page.

THING	HOW MUCH TO TRUST IT
The ~80 employers	Hand researched. Real companies with real early-career hiring.
Openings and links	Checked against each company's hiring system on the date in the footer. Jobs close, so recheck.
Timing status	Verified on that same date.
Scores and explanations	Claude's judgment about you. Useful, arguable, not fact.
Suggested employers	Unverified leads. Confirm each one yourself.

A tailoring run does not re-check the job market. It re-scores and re-explains the existing research. It cannot tell you a job closed yesterday. To actually refresh the openings, run `refresh.py`, which queries every company's hiring system directly.

8 Where your data goes

Short version: your resume goes to Anthropic once and nowhere else, and nothing reaches this project's author.

- **The resume is never stored.** It is read in your browser, included in one request, and discarded. Only Claude's answer is kept.
- **Your API key** is stored in your browser and sent only to Anthropic. This project has no server that could receive it.
- **Saved profiles** stay in that one browser. They do not sync, and clearing site data deletes them.
- **Each person pays their own way.** Runs bill the key that was pasted in, so sharing the link costs you nothing.

The tradeoff for having no server is that your key sits in your browser. Only paste it on a machine you trust, and remember that anyone using your computer could open the page.

9 When something goes wrong

Errors appear in the panel rather than the console, and each one names its own fix.

WHAT YOU SEE	WHAT TO DO
Rejected that API key	The key is wrong or was revoked. Make a new one.
Not allowed to use this model	That account cannot reach that model. Try Sonnet 5.
Rate limited	Wait a minute and run it again.
Could not reach Anthropic	Usually a VPN, an ad blocker or an offline connection.
Ran out of output room	Thinking used the budget. Rerun at a lower effort.
Not valid JSON	A rare bad reply. Just run it again.
Only returned N of 60	Claude stopped early. Run it again.
Could not be saved	Browser storage is full. Delete an old profile.
Nothing happens at all	Serve the folder over http rather than opening the file directly.

If the button itself is missing or dead, the most likely cause is an extension blocking the code library the page loads on first use. The full technical write-up lives in the [project README](#).